

INTRODUCTION



TO STATISTICAL LEARNING

Introduction to Statistical Learning

Support Vector Machines - Class 11

Giora Simchoni

gsimchoni@gmail.com and add #intro2s1 in subject

Stat. and OR Department, TAU

INTRODUCTION



TO STATISTICAL LEARNING

Maximum margin classifier

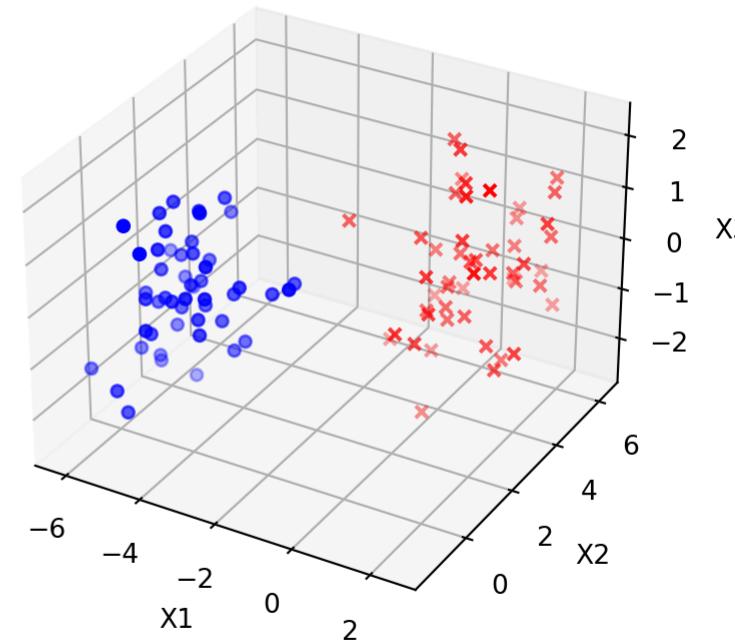
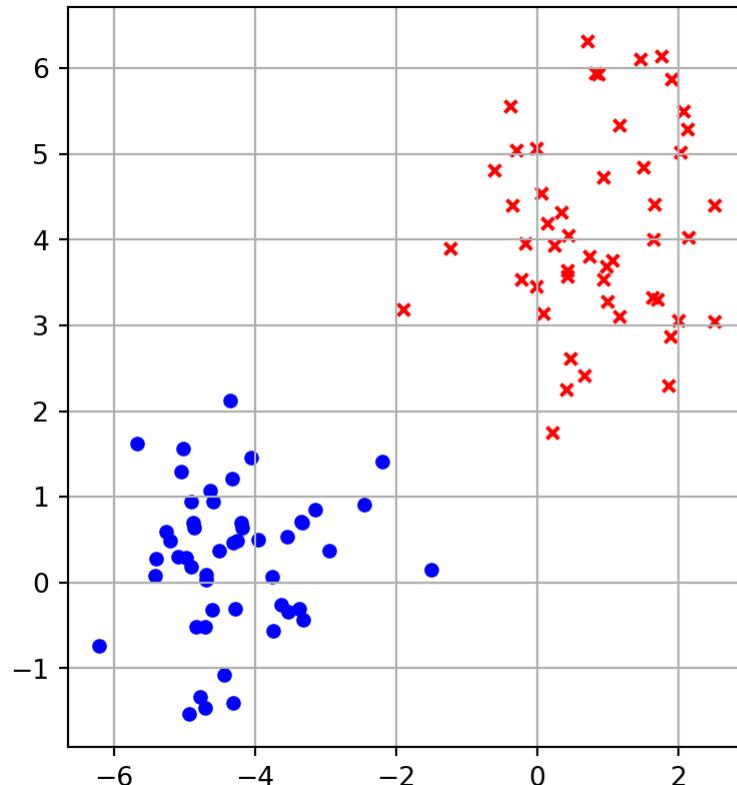
INTRODUCTION



TO STATISTICAL LEARNING

A big assumption

- Suppose $y \in \{-1, 1\}$
- Suppose the classes can be separated by a hyperplane, e.g. for data $T = \{(x_1, y_1), \dots, (x_n, y_n)\}$:

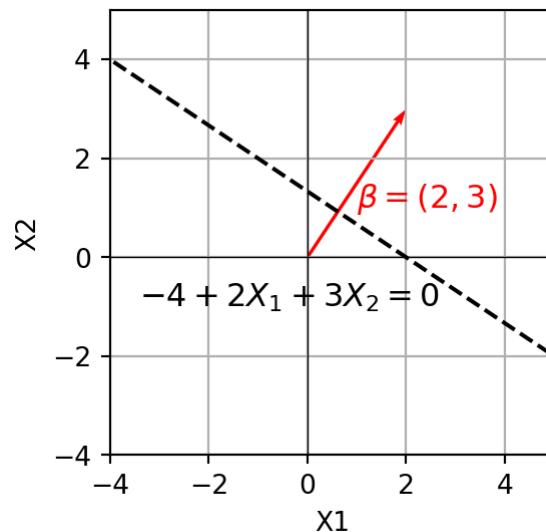


Separating hyperplanes (I)

- $\beta_0, \beta_1, \dots, \beta_p$ define a p -dimensional plane for all points $x \in \mathbb{R}^p$ satisfying:

$$\beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p = \beta_0 + x^T \beta = 0$$

- We say $\beta^* = (\beta_1, \dots, \beta_p) / \|\beta\|$ is normal to the hyperplane



What would be a natural decision rule for separating $y \in \{-1, 1\}$?

Separating hyperplanes (II)

- In other words, a hyperplane is “separating” iff:

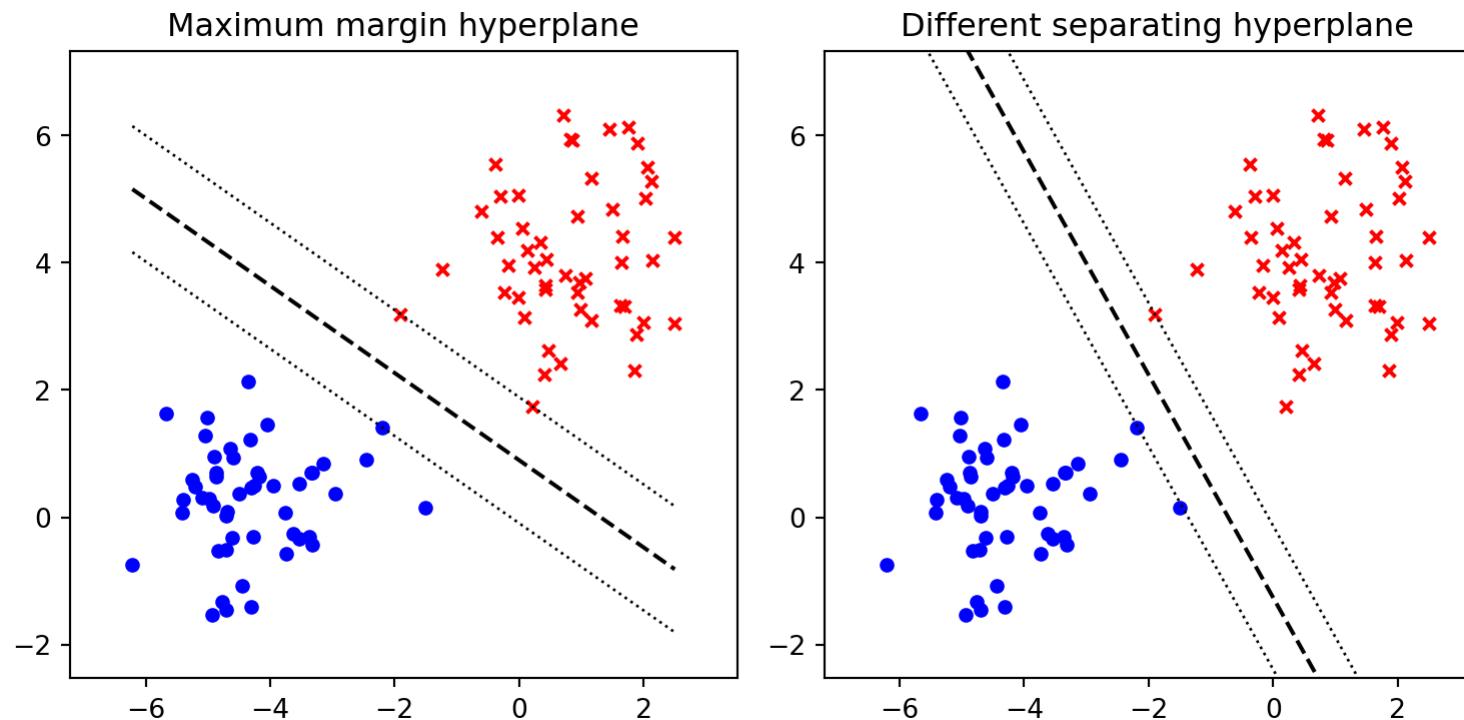
$$y_i \cdot (\beta_0 + x_i^T \beta) > 0 \quad \forall i$$

- And a natural decision rule for separating hyperplanes:

$$\hat{f}(x_0) = \text{sign} \left[\hat{\beta}_0 + x_0^T \hat{\beta} \right]$$

Maximum margin classifier

- Intuitively, we would like the hyperplane that is “fartherst” from points on both sides
- That is, the hyperplane that maximizes the **margin**:
 - the minimum distance of training points to the hyperplane



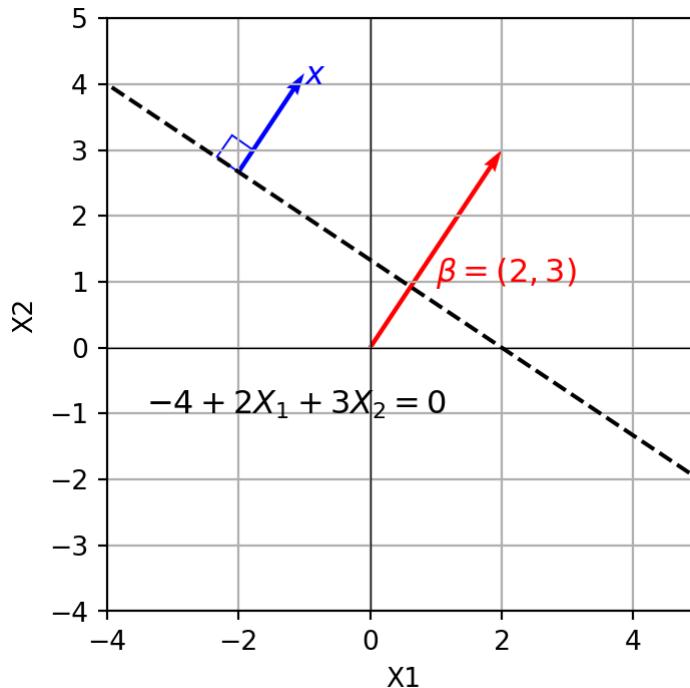
Maximum margin criterion

- So far we have:

$$\max_{\beta, \beta_0} \text{"minimum distance"} \quad s.t. \quad y_i(\beta_0 + x_i^T \beta) > 0 \quad \forall i$$

- The distance between any point x to the hyperplane is: $\frac{|\beta_0 + x^T \beta|}{\|\beta\|_2}$

- because this distance is the length of a vector proportional to the normal β :



Maximum margin criterion

- So for any $M > 0$ a compact way of writing the criterion is:

$$\max_{\beta, \beta_0} M \quad s.t. \quad \|\beta\| = 1 \text{ and } y_i(\beta_0 + x_i^T \beta) \geq M \quad \forall i$$

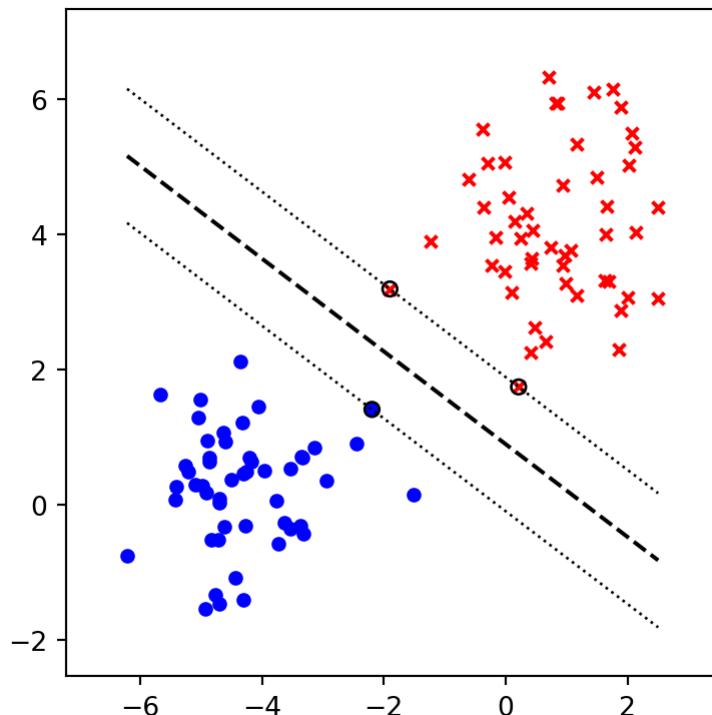
- All the points are at least a distance M from the decision boundary defined by β, β_0 , and we seek β, β_0 that get the largest such M
- Equivalently if we insist on setting $\|\beta\| = 1/M$ we can write:

$$\min_{\beta, \beta_0} \|\beta\| \quad s.t. \quad y_i(\beta_0 + x_i^T \beta) \geq 1 \quad \forall i$$

- because $M = 1/\|\beta\|$ and maximizing M is minimizing $\|\beta\|$
- This is a convex optimization problem (quadratic criterion if we write $\|\beta\|^2$, linear inequality constraints), many efficient solvers exist

Support vectors

Back to the maximum margin classifier:



- The final classifier depends only on support vectors but was reached given all the training data
- Advantages/disadvantages to a classifier that only depends on few observations
- What would logistic regression do?
- What if there is no separable hyperplane?
- And as usual: this is very specific to binary classification, can we generalize?

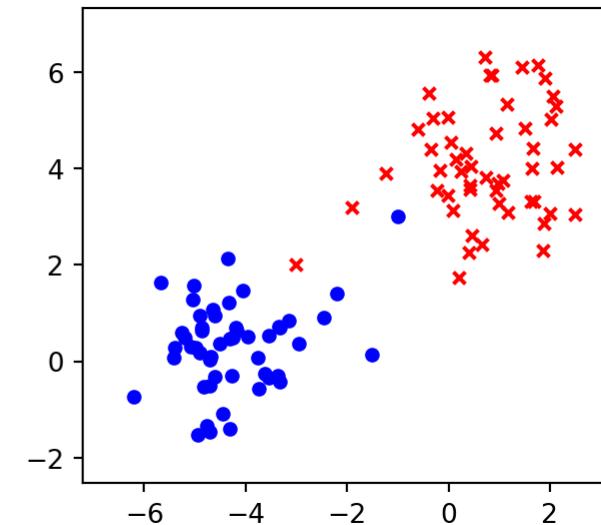
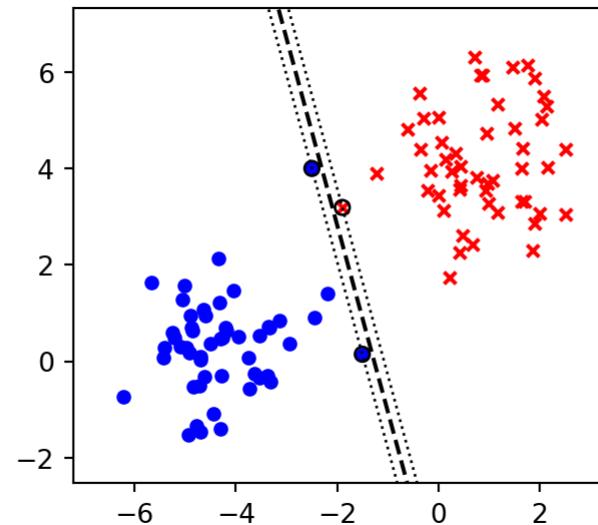
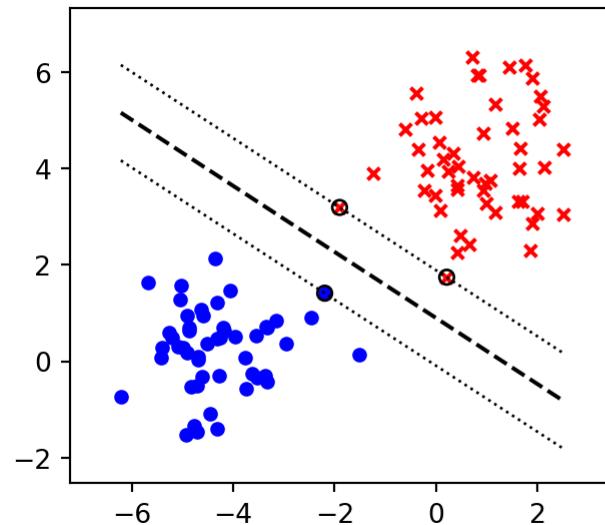
Support vector classifier

INTRODUCTION



TO STATISTICAL LEARNING

Not just the non-separable case



Support vectors classifier

- Let us relax our constraints:
 - Allow observations to be inside the margin
 - Or even on the wrong side of the hyperplane!
 - With a “budget” C for these violations (or penalty)
- We do this with the help of “slack variables”: $\epsilon = (\epsilon_1, \dots, \epsilon_n)$
- The new optimization problem:

$$\max_{\beta, \beta_0, \epsilon} M \quad s.t. \quad \|\beta\| = 1, \quad y_i(\beta_0 + x_i^T \beta) \geq M(1 - \epsilon_i), \quad \epsilon_i \geq 0, \quad \sum_{i=1}^n \epsilon_i \leq C$$

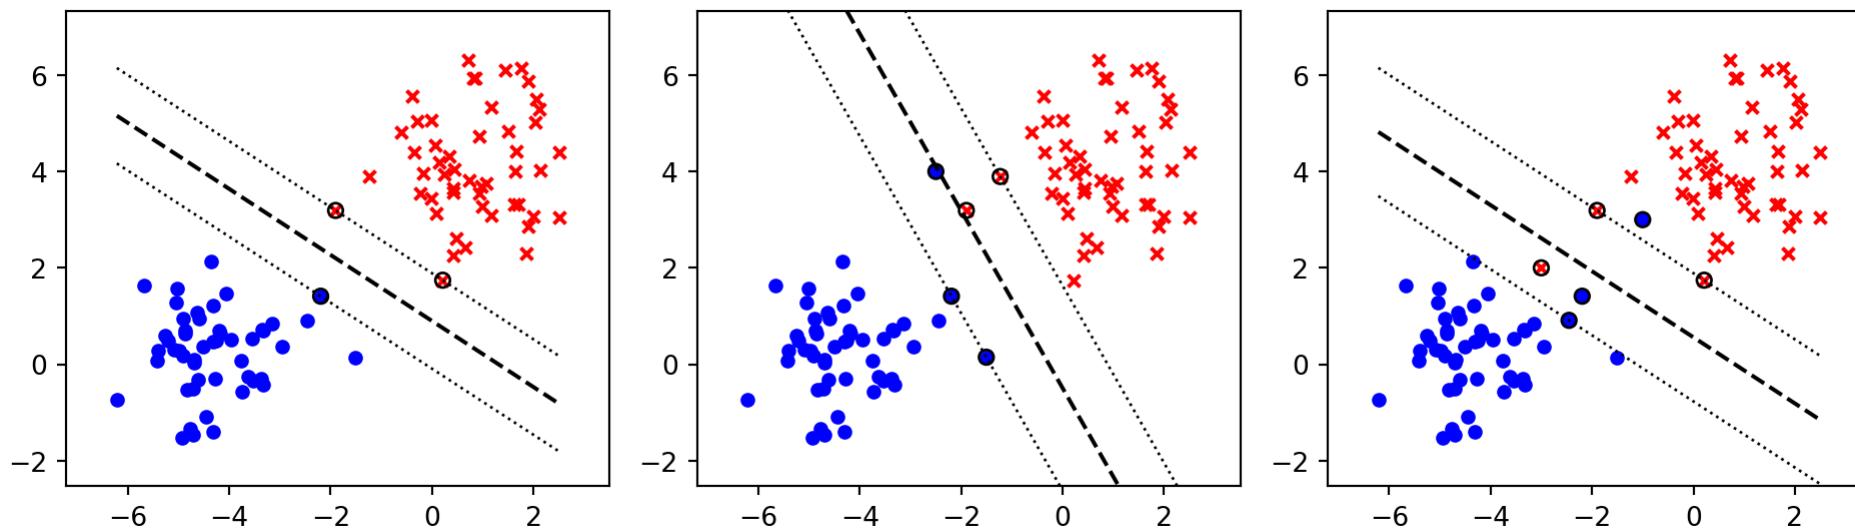
- The final prediction is still:

$$\hat{f}(x_0) = \text{sign} \left[\hat{\beta}_0 + x_0^T \hat{\beta} \right]$$

SVC: slack variables

$$\max_{\beta, \beta_0, \epsilon} M \quad s.t. \quad \|\beta\| = 1, \quad y_i(\beta_0 + x_i^T \beta) \geq M(1 - \epsilon_i), \quad \epsilon_i \geq 0, \quad \sum_{i=1}^n \epsilon_i \leq C$$

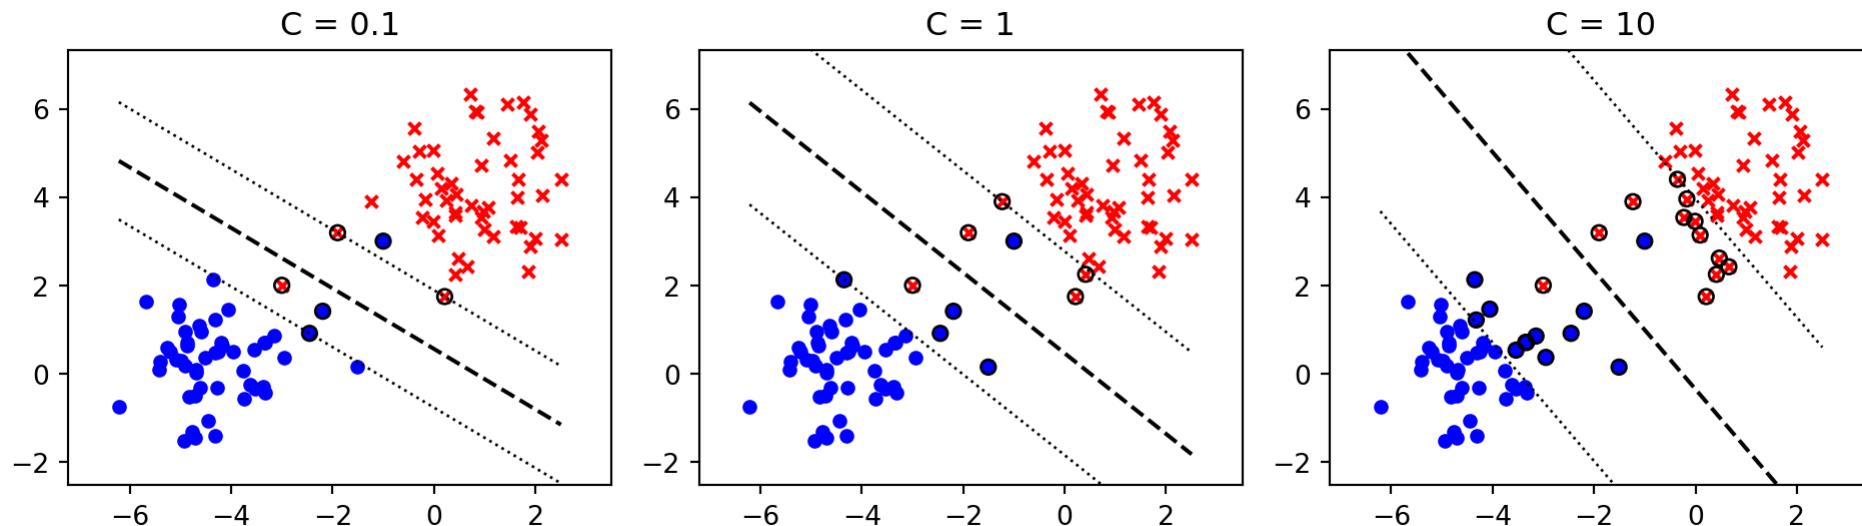
- ϵ_i is the amount by which observation x_i “violates” the margin:
 - If $\epsilon_i = 0$ then x_i is on the correct side of the margin
 - If $\epsilon_i > 0$ then x_i is on the wrong side of the margin
 - If $\epsilon_i > 1$ then x_i is on the wrong side of the hyperplane!



SVC: the C parameter

$$\max_{\beta, \beta_0, \epsilon} M \quad s.t. \quad \|\beta\| = 1, \quad y_i(\beta_0 + x_i^T \beta) \geq M(1 - \epsilon_i), \quad \epsilon_i \geq 0, \quad \sum_{i=1}^n \epsilon_i \leq C$$

- C is the amount by which the margin may be “violated”:
 - If $C = 0$ then $\epsilon = \mathbf{0}$, back to maximum margin classifier
 - If $C > 0$ no more than C observations can be on the wrong side of the hyperplane



How to choose C ? What is the relation between C and the bias-variance tradeoff?

SVC: equivalent forms

- As with maximum margin classifier, if we insist on setting $\|\beta\| = 1/M$ we can write:

$$\min_{\beta, \beta_0, \epsilon} \|\beta\| \quad s.t. \quad y_i(\beta_0 + x_i^T \beta) \geq 1 - \epsilon_i, \quad \epsilon_i \geq 0, \quad \sum_{i=1}^n \epsilon_i \leq C \quad \forall i$$

- In fact, it is more common to see the equivalent form:

$$\min_{\beta, \beta_0, \epsilon} \frac{1}{2} \|\beta\|^2 + C \sum_{i=1}^n \epsilon_i \quad s.t. \quad y_i(\beta_0 + x_i^T \beta) \geq 1 - \epsilon_i, \quad \epsilon_i \geq 0 \quad \forall i$$

- In which case, notice the C hyperparameter is now a *penalty* (this is also more similar to how sklearn sees it)

Support vector machines (SVM)

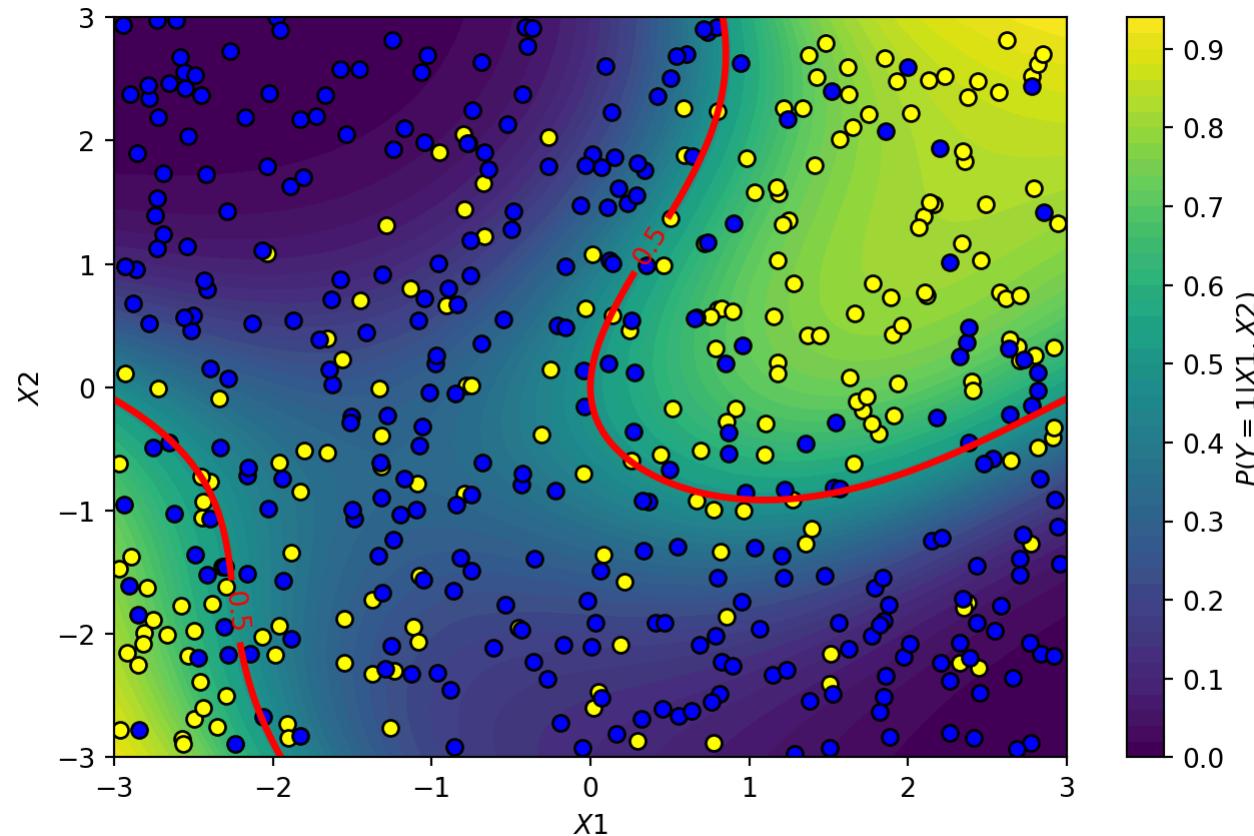
INTRODUCTION



TO STATISTICAL LEARNING

The non-linear case

Recall the Bayes decision boundary example from our Bias-Variance discussion:



Unfortunately, “real data” looks more like this.

Adding polynomial terms?

- Similar to polynomial regression we could add quadratic terms:

$$\max_{\beta^1, \beta^2, \beta_0, \epsilon} M \quad s.t. \|\beta^1\|^2 + \|\beta^2\|^2 = 1, \quad y_i(\beta_0 + x_i^T \beta^1 + (x_i^2)^T \beta^2) \geq M(1 - \epsilon_i)$$

$$\epsilon_i \geq 0, \quad \sum_{i=1}^n \epsilon_i \leq C$$

- What about cubic terms? What about interactions?
- Like in regression we could *expand* x_i to any feature mapping $h(x_i) : \mathbb{R}^p \rightarrow \mathbb{R}^q$ and continue as is
- But this becomes really high-dimensional, really fast.

The kernel trick (I)

- An amazing insight, after a lot of algebra, our solution can be written as:

$$f(x_0) = \text{sign} [\beta_0 + x_0^T \beta] = \text{sign} \left[\beta_0 + \sum_{i=1}^n \alpha_i \langle x_i, x_0 \rangle \right]$$

- where:
 - $\langle x_i, x_0 \rangle$ is the inner product $x_i^T x_0 = \sum_{j=1}^p x_{ij} x_{0j}$
 - $\alpha_i, \dots, \alpha_n$ are n parameters for n observations,
 - but $\alpha_i > 0$ only for support vectors, otherwise $\alpha_i = 0$
- So for any $h(x_i)$, the solution can be written in terms of inner products only:

$$f(x_0) = \text{sign} [\beta_0 + h(x_0)^T \beta] = \text{sign} \left[\beta_0 + \sum_{i=1}^n \alpha_i \langle h(x_i), h(x_0) \rangle \right]$$

The kernel trick (II)

- Now suppose $x_i = (x_{i1}, x_{i2})$ and $h : \mathbb{R}^2 \rightarrow \mathbb{R}^6$ is:

$$h(x_i) = (1, \sqrt{2}x_{i1}, \sqrt{2}x_{i2}, x_{i1}^2, x_{i2}^2, \sqrt{2}x_{i1}x_{i2})$$

- Amazingly, the inner product with any point x_0 is quite compact:

$$\begin{aligned} \langle h(x_i), h(x_0) \rangle &= (1 \quad \sqrt{2}x_{i1} \quad \sqrt{2}x_{i2} \quad x_{i1}^2 \quad x_{i2}^2 \quad \sqrt{2}x_{i1}x_{i2}) \\ &\quad \left(\begin{array}{c} 1 \\ \sqrt{2}x_{01} \\ \sqrt{2}x_{02} \\ x_{01}^2 \\ x_{02}^2 \\ \sqrt{2}x_{01}x_{02} \end{array} \right) = \\ &= (1 + x_i^T x_0)^2 \end{aligned}$$

 Why is that helpful?

The kernel trick (III)

- Let this be a **kernel** function:

$$\langle h(x_i), h(x_0) \rangle = (1 + x_i^T x_0)^2 = K(x_i, x_0)$$

- Back to our solution:

$$f(x_0) = \text{sign} \left[\beta_0 + \sum_{i=1}^n \alpha_i \langle h(x_i), h(x_0) \rangle \right] = \text{sign} \left[\beta_0 + \sum_{i=1}^n \alpha_i K(x_i, x_0) \right]$$

- The kernel trick:**
 - Forget about specifying $h(x_i) : \mathbb{R}^p \rightarrow \mathbb{R}^q$!
 - Focus on specifying kernel functions $K(x_i, x_j) : \mathbb{R}^p \times \mathbb{R}^p \rightarrow \mathbb{R}$
 - That is, flexible similarity functions between feature vectors

Common kernels

kernel	$K(x_i, x_j)$	comment
Linear	$\langle x_i, x_j \rangle$	SVC!
Polynomial	$(1 + \langle x_i, x_j \rangle)^d$	For $p = 2, d = 2$ we got $h(x)$
Gaussian/SE/RBF	$\exp(-\gamma \ x_i - x_j\ ^2)$	infinite feature mapping space!
Sigmoid	$\tanh(\kappa_1 \langle x_i, x_j \rangle + \kappa_2)$	

- Each of these *implicitly* uses some mapping $h(x)$, but always:

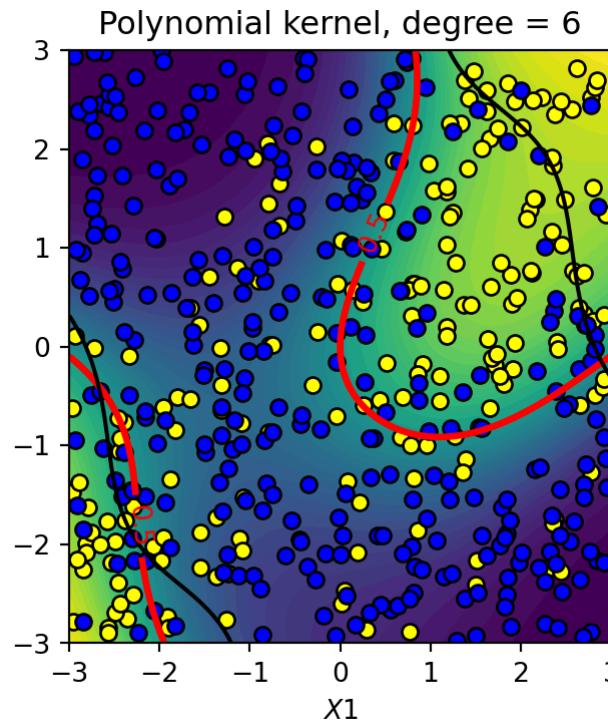
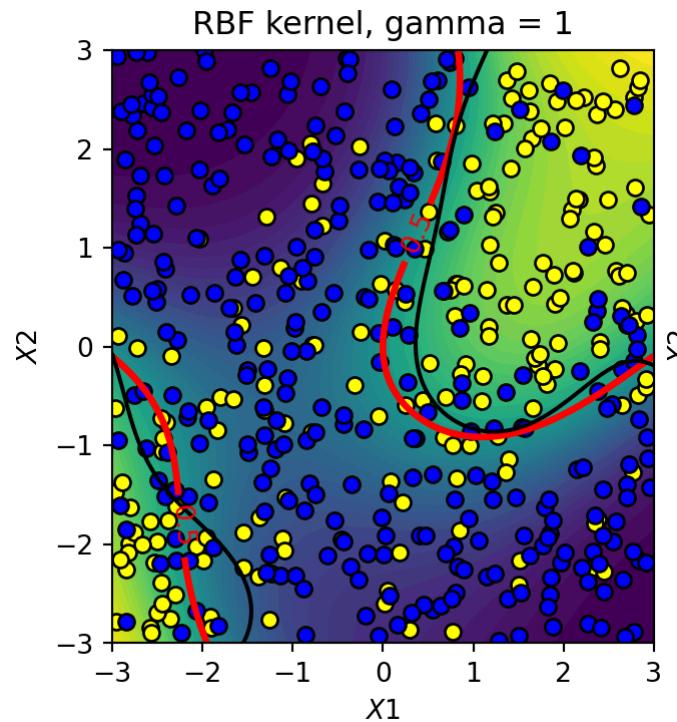
$$K(x_i, x_j) = \langle h(x_i), h(x_j) \rangle$$
- On entire training set need to compute: $K_{n \times n}$ (kernel matrix) where

$$K_{ij} = K(x_i, x_j) = \langle h(x_i), h(x_j) \rangle$$



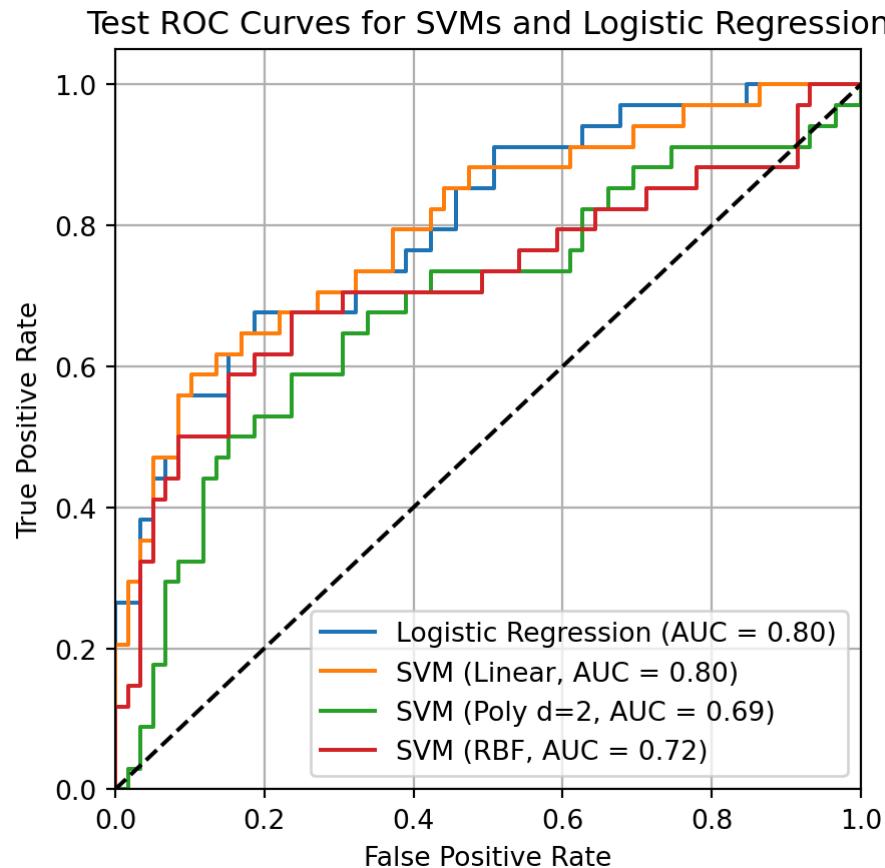
Can any $K(x_i, x_j) : \mathbb{R}^p \times \mathbb{R}^p \rightarrow \mathbb{R}$ function use as a kernel?

Support vector machines (SVM)



Notice this might add a few hyperparameters, including the kernel itself!

Example: SAHeart data



Notice to compute ROC SVM needs to output a probability or score.

This is out of scope, but you could think of the distance of an observation to the hyperplane as a type of confidence.

SVM extensions

INTRODUCTION

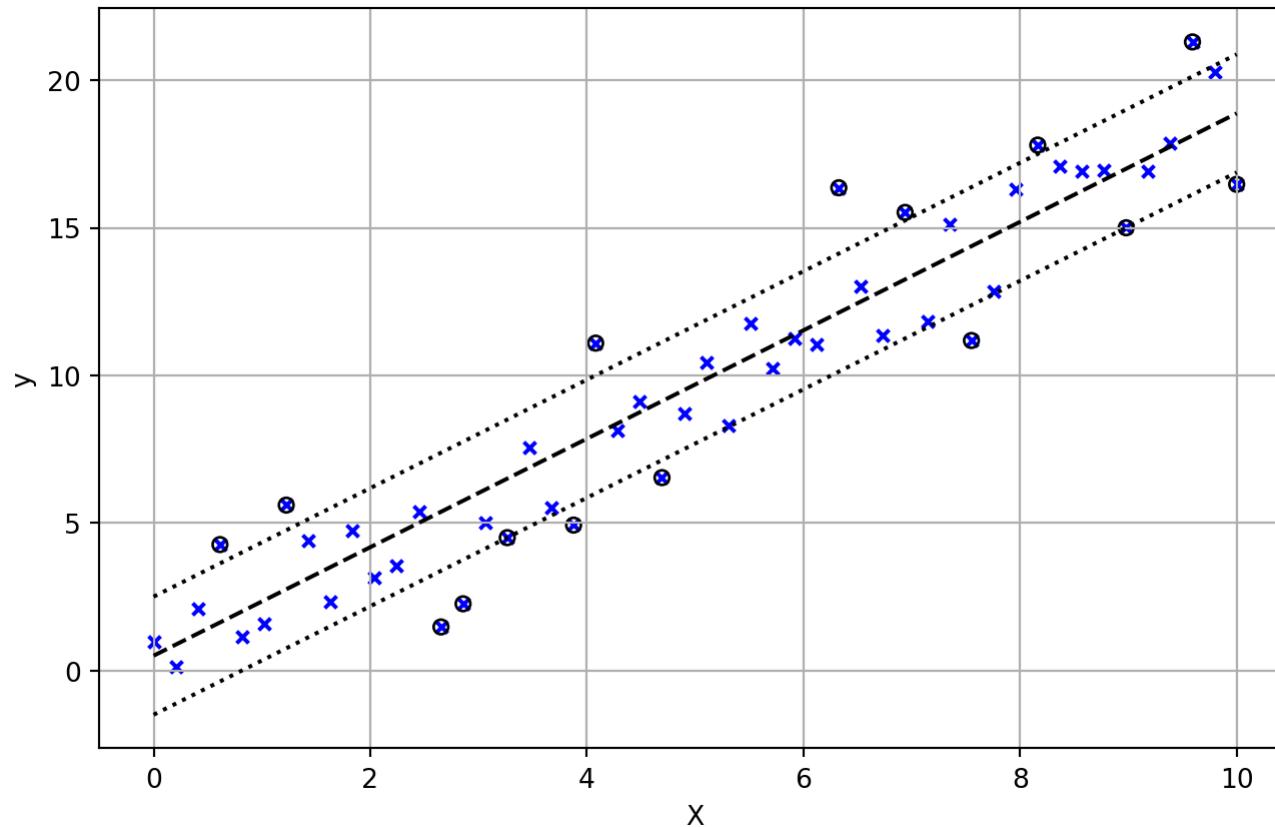


TO STATISTICAL LEARNING

SVM for K classes

- One-versus-one (OVO):
 - Run all $\binom{K}{2}$ pairwise models: class k vs. class k'
 - Assign x_0 to class k to which it was most frequently assigned to
- One-versus-rest (OVR):
 - Run all K models: class k (+1) vs. remaining $K - 1$ classes (-1)
 - Assign x_0 to class k in which it gets the highest distance (confidence) from hyperplane $\beta_{0k} + x_0^T \hat{\beta}_k$

Support vector regression (SVR)



What would be the support vectors for SVR?

SVR Criterion

- From maximum margin classifier:

$$\max_{\beta, \beta_0} M \quad s.t. \quad \|\beta\| = 1 \text{ and } y_i(\beta_0 + x_i^T \beta) \geq M \quad \forall i$$

- To “maximum margin regressor” (not really a thing):

$$\max_{\beta, \beta_0} M \quad s.t. \quad \|\beta\| = 1 \text{ and } |y_i - (\beta_0 + x_i^T \beta)| \leq M \quad \forall i$$

- From SVC:

$$\max_{\beta, \beta_0, \epsilon} M \quad s.t. \quad \|\beta\| = 1, \quad y_i(\beta_0 + x_i^T \beta) \geq M(1 - \epsilon_i), \quad \epsilon_i \geq 0, \quad \sum_{i=1}^n \epsilon_i \leq C$$

- To SVR:

$$\max_{\beta, \beta_0, \epsilon} M \quad s.t. \quad \|\beta\| = 1, \quad |y_i - (\beta_0 + x_i^T \beta)| \leq M + \epsilon_i, \quad \epsilon_i \geq 0, \quad \sum_{i=1}^n \epsilon_i \leq C$$

- Though this is more common:

$$\min_{\beta, \beta_0, \epsilon} \frac{1}{2} \|\beta\|^2 + C \sum_{i=1}^n \epsilon_i \quad s.t. \quad |y_i - (\beta_0 + x_i^T \beta)| \leq M + \epsilon_i, \quad \epsilon_i \geq 0$$

- And of course SVR can be kernelized as well