

INTRODUCTION



TO STATISTICAL LEARNING

Introduction to Statistical Learning

Generative Models - Class 12

Giora Simchoni

gsimchoni@gmail.com and add #intro2s1 in subject

Stat. and OR Department, TAU

INTRODUCTION



TO STATISTICAL LEARNING

Intro. to Generative Models

INTRODUCTION



TO STATISTICAL LEARNING

Generated images



Face recognition using eigenfaces by Turk & Pentland (1991)



Image generated using DALL·E by OpenAI with: “a realistic image of a dog and a cat hugging in front of the eiffel tower”

How did we get here?

Discriminative models

- We are constantly trying to *model* $E(Y|X = x)$ (regression) or $P(Y|X = x)$ (Bayes classifier):
 - Linear:
 - linear regression, ridge, lasso
 - logistic regression
 - Non-linear:
 - KNN
 - Trees, RF, Boosting



Even our Bias-variance tradeoff analysis assumed a Fixed- X scenario!

Generative models

But for classification:

$$P(Y|X) = \frac{P(X, Y)}{P(X)}$$

- **Generative** models focus on modeling the joint distribution $P(X, Y)$, or more specifically:

$$P(Y|X) = \frac{P(X, Y)}{P(X)} = \frac{P(Y)P(X|Y)}{P(X)} \propto P(Y)P(X|Y)$$

- Focus on the mechanism which **generated** the data, not just the data
- Especially useful when n is small

Classification using generative models

- Even more specifically, if $Y \in \{1, \dots, K\}$:

$$P(Y = k|X = x) = \frac{P(Y = k)P(X = x|Y = k)}{P(X = x)} = \frac{\pi_k f_k(x)}{\sum_{l=1}^K \pi_l f_l(x)}$$

- We focus on estimating $\pi_k, f_k(x)$



And if we are so good at estimating $f_k(x) = P(X = x|Y = k)$ why not generate more!

Detour: multivariate normal

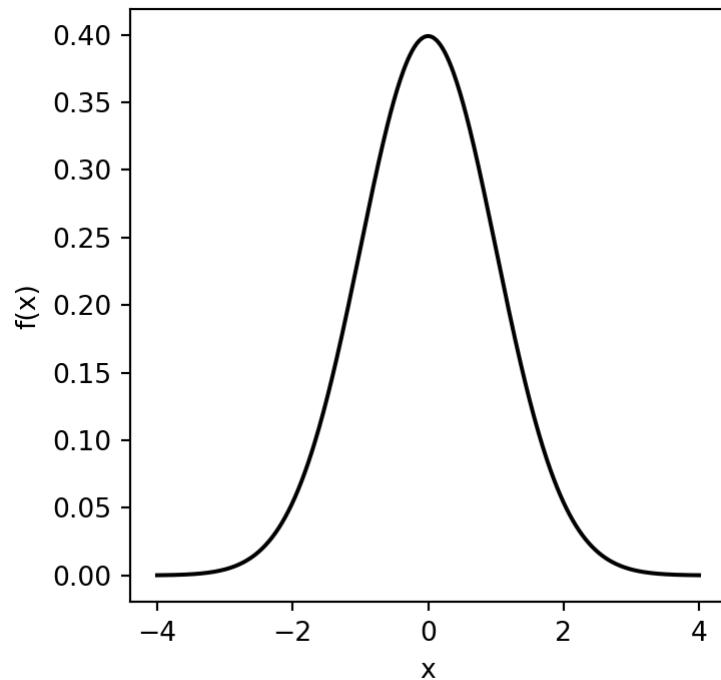
INTRODUCTION



TO STATISTICAL LEARNING

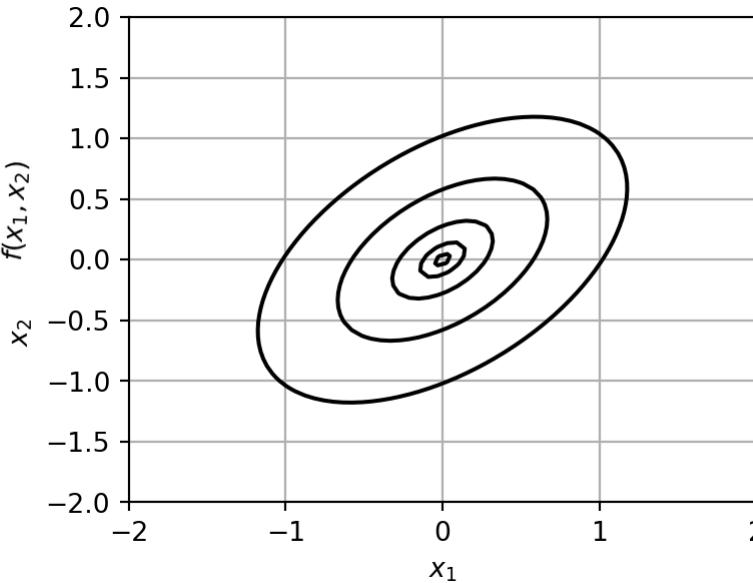
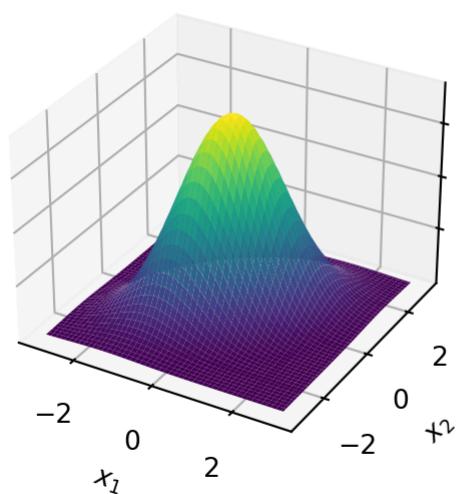
Univariate normal

- $X \in \mathbb{R} \sim \mathcal{N}(\mu, \sigma^2) \Rightarrow E(X) = \mu; \quad V(X) = \sigma^2 > 0$
- $f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2\sigma^2}(x - \mu)^2\right)$
- $F(x) = P(X \leq x) = \int_{-\infty}^x f(x)dx = \Phi\left(\frac{x-\mu}{\sigma}\right)$



Bivariate normal

- $X \in \mathbb{R}^2 \sim \mathcal{BVN}(\mu, \Sigma) \Rightarrow E(X) = \mu = \begin{pmatrix} \mu_1 \\ \mu_2 \end{pmatrix}; \quad V(X) = \Sigma = \begin{pmatrix} \sigma_1^2 & \rho\sigma_1\sigma_2 \\ \rho\sigma_1\sigma_2 & \sigma_2^2 \end{pmatrix}$
- $\Leftrightarrow X_1 \sim \mathcal{N}(\mu_1, \sigma_1^2), X_2 \sim \mathcal{N}(\mu_2, \sigma_2^2); \quad \rho = \text{Cor}(X_1, X_2) = \frac{\text{Cov}(X_1, X_2)}{\sigma_1\sigma_2}$
- $f(x) = f(x_1, x_2) = \frac{1}{2\pi\sigma_1\sigma_2\sqrt{1-\rho^2}} \exp\left(-\frac{1}{2(1-\rho^2)} \left[\frac{(x_1-\mu_1)^2}{\sigma_1^2} + \frac{(x_2-\mu_2)^2}{\sigma_2^2} - \frac{2\rho(x_1-\mu_1)(x_2-\mu_2)}{\sigma_1\sigma_2} \right]\right)$
- $F(x) = P(X_1 \leq x_1, X_2 \leq x_2) = \int_{-\infty}^{x_2} \int_{-\infty}^{x_1} f(x_1, x_2) dx_1 dx_2$



Multivariate normal

- $X \in \mathbb{R}^p \sim \mathcal{MVN}(\mu, \Sigma)$

- $E(X) = \mu = \begin{pmatrix} \mu_1 \\ \vdots \\ \mu_p \end{pmatrix}; \quad V(X) = \Sigma = \begin{pmatrix} \sigma_1^2 & \cdots & \rho_{1,p}\sigma_1\sigma_p \\ \vdots & \ddots & \vdots \\ \rho_{1,p}\sigma_1\sigma_p & \cdots & \sigma_p^2 \end{pmatrix}$

- $\Leftrightarrow X_1 \sim \mathcal{N}(\mu_1, \sigma_1^2), \dots, X_p \sim \mathcal{N}(\mu_p, \sigma_p^2); \quad \rho_{j,k} = \text{Cor}(X_j, X_k) = \frac{\text{Cov}(X_j, X_k)}{\sigma_j \sigma_k}$
- $f(x) = f(x_1, \dots, x_p) = \frac{1}{(2\pi)^{p/2} |\Sigma|^{1/2}} \exp \left(-\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right)$



How many unique params in Σ ?

Linear Discriminant Analysis

INTRODUCTION



TO STATISTICAL LEARNING

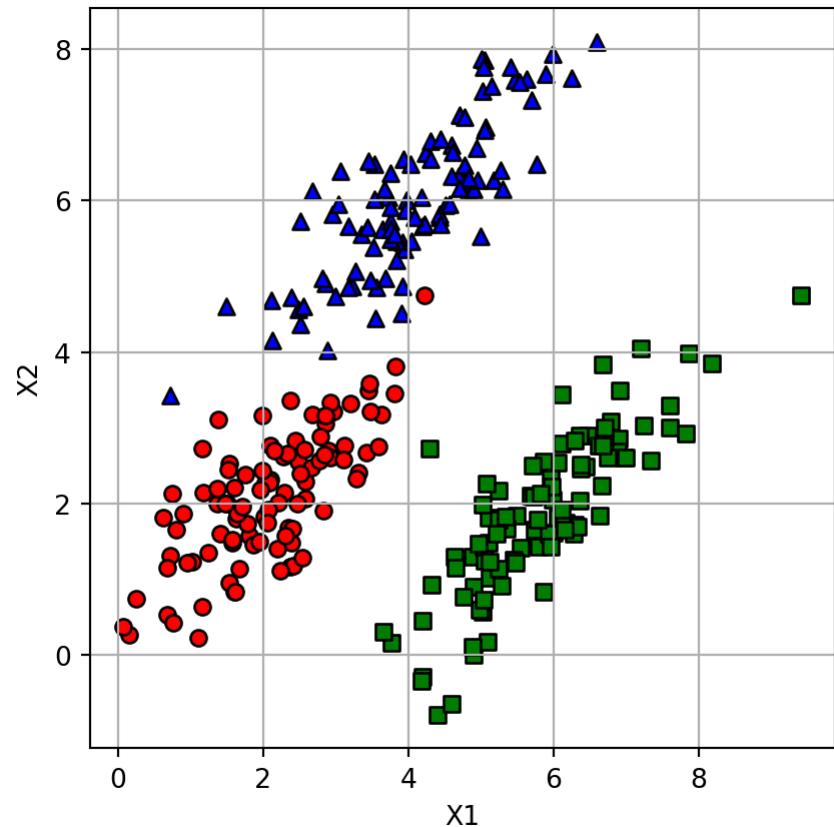
Linear discriminant analysis (LDA)

- Recall:

$$P(Y = k|X = x) = \frac{P(Y = k)P(X = x|Y = k)}{P(X = x)} = \frac{\pi_k f_k(x)}{\sum_{l=1}^K \pi_l f_l(x)}$$

- $\hat{f}(x_0) = \arg \max_k \pi_k f_k(x_0)$
- In LDA:
 - π_k are priors, or (spoiler): $\hat{\pi}_k = \frac{\sum_{i=1}^n \mathbb{I}[Y_i=k]}{n}$
 - $f_k(x)$ are multivariate Gaussian, or: $X|Y = k \sim \mathcal{N}(\mu_k, \Sigma)$
 - Notice the covariance matrix Σ is the same $\forall k$

LDA: What to expect



What do you expect the decision rule(s) to look like?

LDA: $k - j$ classes decision rule

- Why choose class k over j ?

$$P(Y = k|X = x) > P(Y = j|X = x)$$

$$\Leftrightarrow \pi_k f_k(x) > \pi_j f_j(x)$$

$$\Leftrightarrow \log [\pi_k f_k(x)] > \log [\pi_j f_j(x)]$$

- Assume $\pi_k, \pi_j, \mu_k, \mu_j, \Sigma$ known
- $X|Y = k \sim \mathcal{N}(\mu_k, \Sigma) \Rightarrow f_k(x) = \frac{1}{(2\pi)^{p/2}|\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x - \mu_k)^T \Sigma^{-1}(x - \mu_k)\right)$
- $\log [\pi_k f_k(x)] = \log(\pi_k) + \log(f_k(x)) = \log(\pi_k) + C - \frac{1}{2}(x - \mu_k)^T \Sigma^{-1}(x - \mu_k)$

LDA: $k - j$ classes decision rule

So, select class k over class j if $\log [\pi_k f_k(x)] > \log [\pi_j f_j(x)]$ means:

$$\log(\pi_k) - \frac{1}{2}(x - \mu_k)^T \Sigma^{-1}(x - \mu_k) > \log(\pi_j) - \frac{1}{2}(x - \mu_j)^T \Sigma^{-1}(x - \mu_j)$$

$$\delta_{k>j}(x) : x^T \Sigma^{-1}(\mu_k - \mu_j) + \left[-\frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \frac{1}{2} \mu_j^T \Sigma^{-1} \mu_j + \log(\pi_k) - \log(\pi_j) \right] > 0$$



What shape is $\delta_{k>j}(x)$?

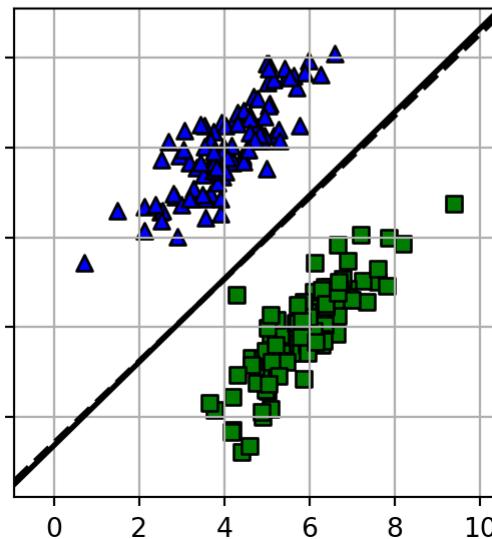
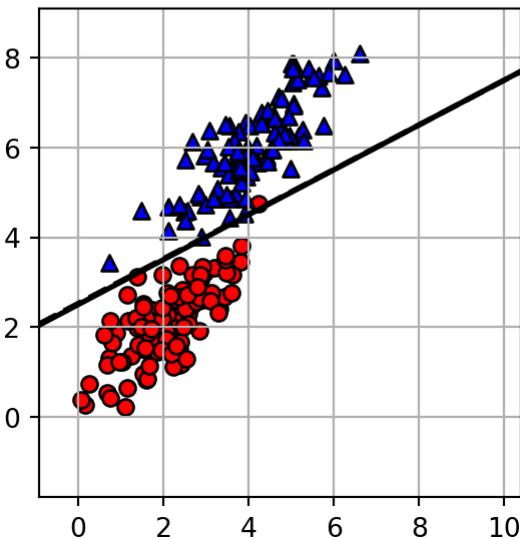
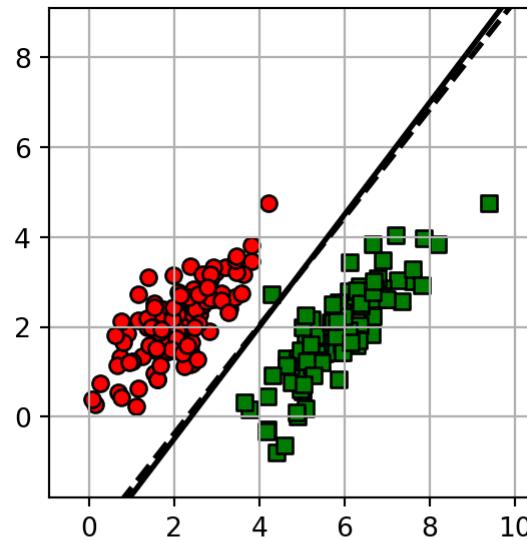
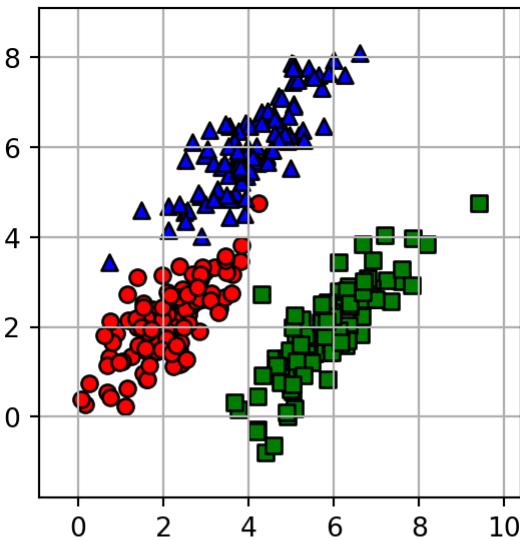
LDA: Estimation

- If we want to fit LDA to data, we need to estimate the parameters:
 $\mu_1, \dots, \mu_K, \Sigma, \pi_1, \dots, \pi_K$.
- This is naturally done from the data (e.g. by maximum likelihood):

$$\hat{\pi}_k = \frac{\sum_{i=1}^n \mathbb{I}[Y_i = k]}{n}, \quad \hat{\mu}_k = \frac{\sum_{i=1}^n \mathbb{I}[Y_i = k] x_i}{\sum_{i=1}^n \mathbb{I}[Y_i = k]}$$

$$\hat{\Sigma} = \frac{1}{n - K} \sum_{k=1}^K \sum_{\mathbb{I}[Y_i = k]} (x_i - \hat{\mu}_k)(x_i - \hat{\mu}_k)^T$$

LDA: Example



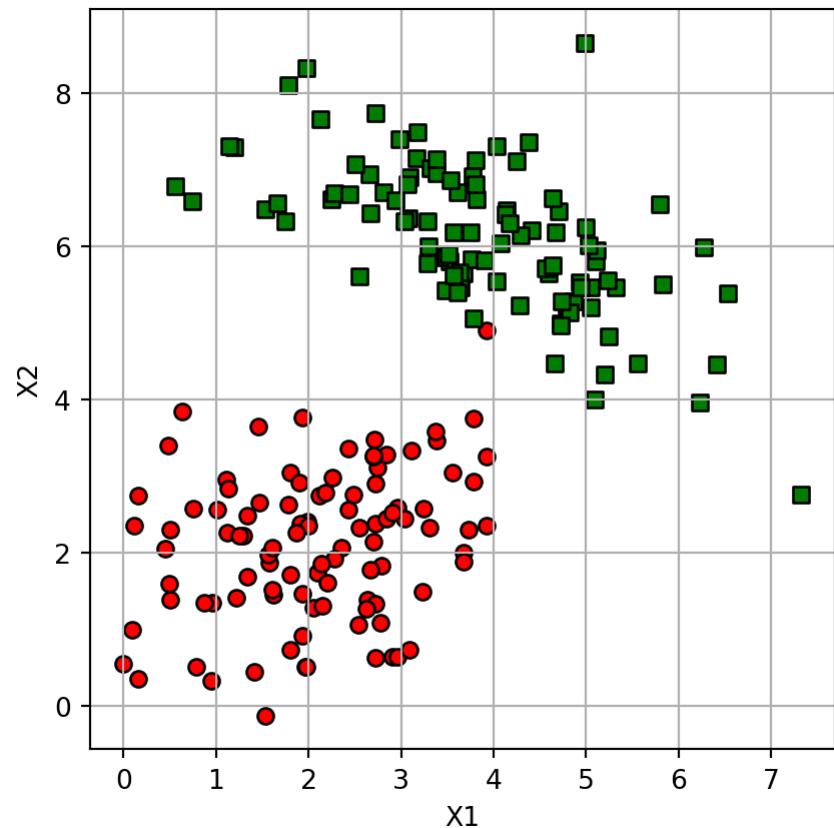
Quadratic Discriminant Analysis

INTRODUCTION



TO STATISTICAL LEARNING

A common covariance?



Quadratic linear analysis (QDA)

- Now assume: $X|Y = k \sim \mathcal{N}(\mu_k, \Sigma_k)$
- Decision rule stays the same:
 - select class k over class j if $\log [\pi_k f_k(x)] > \log [\pi_j f_j(x)]$
- $\log [\pi_k f_k(x)] = \log(\pi_k) + C - \frac{1}{2}(x - \mu_k)^T \Sigma_k^{-1} (x - \mu_k) - \frac{1}{2}\log(|\Sigma_k|)$

$\delta_{k>j}(x) :$

$$\begin{aligned}
 & - \frac{1}{2}x^T (\Sigma_k^{-1} - \Sigma_j^{-1})x + x^T (\Sigma_k^{-1}\mu_k - \Sigma_j^{-1}\mu_j) \\
 & + \left[-\frac{1}{2}\mu_k^T \Sigma^{-1}\mu_k + \frac{1}{2}\mu_j^T \Sigma^{-1}\mu_j + \log(\pi_k) - \log(\pi_j) \right. \\
 & \quad \left. + \frac{1}{2}(\log(|\Sigma_j|) - \log(|\Sigma_k|)) \right] > 0
 \end{aligned}$$

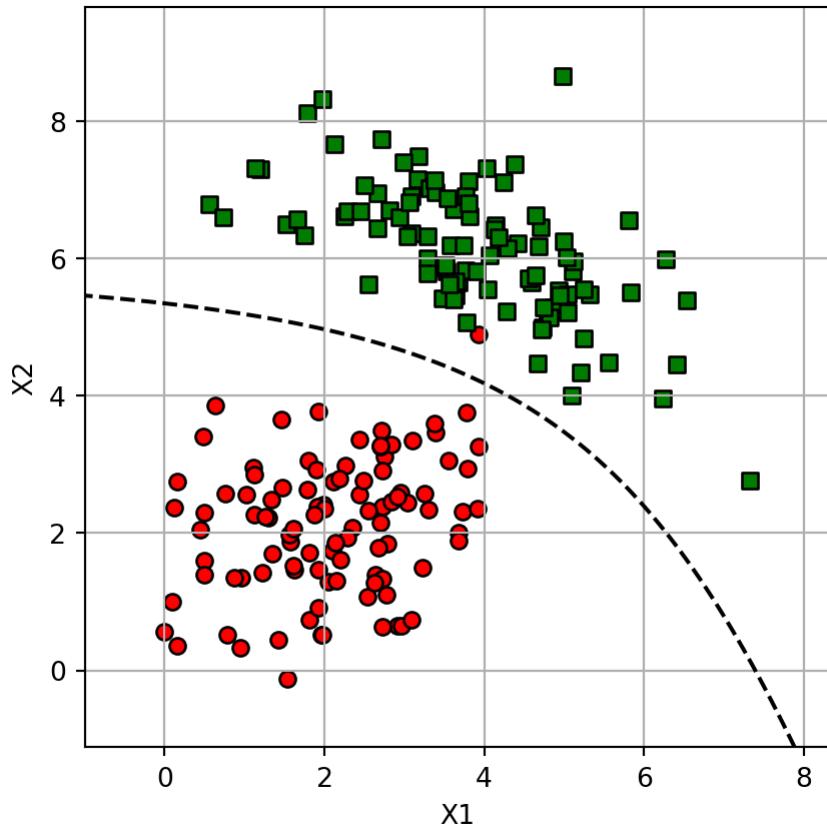


What shape is $\delta_{k>j}(x)$?

QDA: Estimation

- Estimate μ_k, π_k as in LDA and:

- $\hat{\Sigma}_k = \frac{1}{n_k - 1} \sum_{\mathbb{I}[Y_i=k]} (X_i - \hat{\mu}_k)(X_i - \hat{\mu}_k)^T$



How many params are estimated in QDA? LDA?

Naive Bayes

INTRODUCTION



TO STATISTICAL LEARNING

Naive Bayes

- Recall: $P(Y = k|X = x) = \frac{P(Y=k)P(X=x|Y=k)}{P(X=x)} = \frac{\pi_k f_k(x)}{\sum_{l=1}^K \pi_l f_l(x)}$
- In Naive Bayes:
 - π_k are priors, or: $\hat{\pi}_k = \frac{\sum_{i=1}^n \mathbb{I}[Y_i=k]}{n}$
 - Within the k -th class, the p predictors are **independent**
 - $f_k(x) = f_{k1}(x_1) \times f_{k2}(x_2) \times \cdots \times f_{kp}(x_p)$
 - $f_{kj}(x_j)$ for a continuous feature: $\mathcal{N}(\mu_{jk}, \sigma_{jk}^2)$, $Exp(\lambda_{jk})$, KDE, ...
 - $f_{kj}(x_j)$ for a discrete feature: $\hat{f}_{kj}(x_j) = \begin{cases} 0.2 & \text{if } x_j = 1 \\ 0.8 & \text{if } x_j = 2 \end{cases}$

 We don't really believe this independence... when is this naive assumption particularly useful?

 How is Naive Bayes = LDA/QDA?

Comparing classifiers

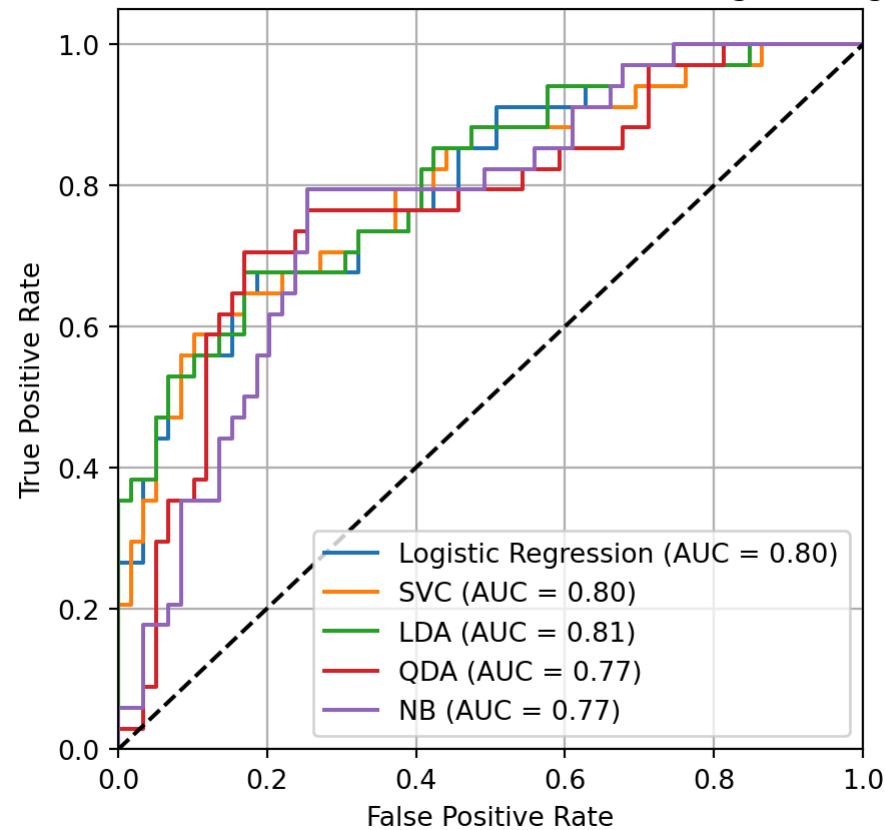
INTRODUCTION



TO STATISTICAL LEARNING

Example: SAHeart data

Test ROC Curves for Generative models and Logistic Regression



LDA vs. logistic regression

- For $K = 2$ (though we can show this for any K), logistic regression:

$$\begin{aligned}\text{logit}(P(Y = 1|X)) &= \log\left(\frac{P(Y = 1|X)}{1 - P(Y = 1|X)}\right) \\ &= \log\left(\frac{P(Y = 1|X)}{P(Y = 0|X)}\right) = \beta_0 + x^T \beta\end{aligned}$$

- For $K = 2$ (though we can show this for any K), LDA:

$$\begin{aligned}\log\left(\frac{P(Y = 1|X)}{P(Y = 0|X)}\right) &= \log\left(\frac{\pi_1 f_1(x)}{\pi_0 f_0(x)}\right) = \log(\pi_1 f_1(x)) - \log(\pi_0 f_0(x)) \\ &= x^T \Sigma^{-1} (\mu_1 - \mu_0) \\ &\quad + \left[-\frac{1}{2} \mu_1^T \Sigma^{-1} \mu_1 + \frac{1}{2} \mu_0^T \Sigma^{-1} \mu_0 + \log(\pi_1) - \log(\pi_0) \right] \\ &= \alpha_0 + x^T \alpha\end{aligned}$$

Similarly for all methods

$$\delta_{1>0}(x) = \log \left(\frac{P(Y=1|X=x)}{P(Y=0|X=x)} \right) =$$

method	type	$\delta_{1>0}(x)$
LR	discriminative, linear	$\beta_0 + x^T \beta$
LDA	generative, linear	$\alpha_0 + x^T \alpha$
QDA	generative, non-linear	$\gamma_0 + x^T \gamma + x^T \Gamma x$
NB	generative, non-linear	$\tau_0 + \sum_{j=1}^p \tau_j(x_j)$